

# Αρχιτεκτονική Υπολογιστών

## Εισαγωγή στο πρόσωμοιωτή Gem5



Αγγελος Αθανασιάδης

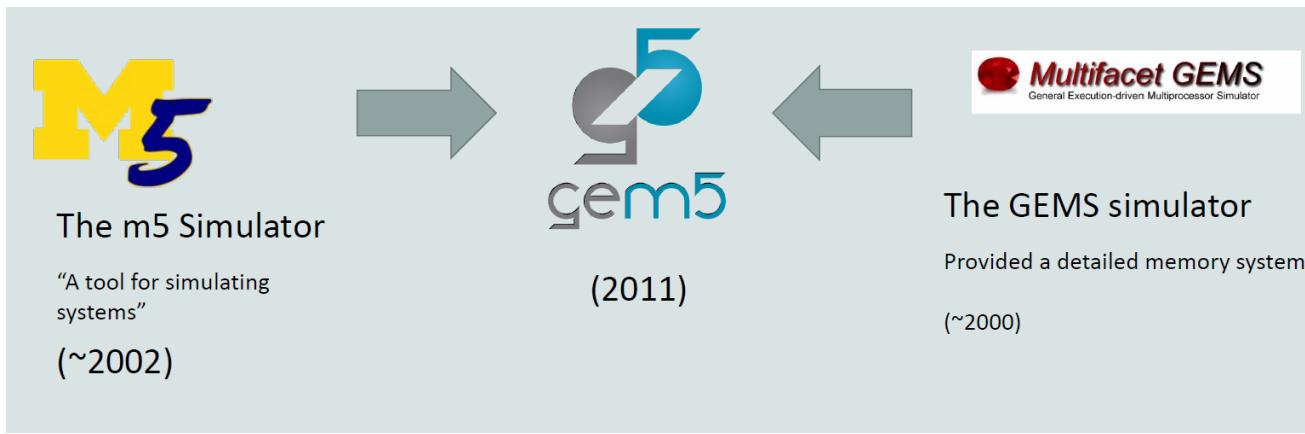


ΤΜΗΜΑ ΗΛΕΚΤΡΟΛΟΓΩΝ ΜΗΧΑΝΙΚΩΝ  
& ΜΗΧΑΝΙΚΩΝ ΥΠΟΛΟΓΙΣΤΩΝ

ΑΡΙΣΤΟΤΕΛΕΙΟ ΠΑΝΕΠΙΣΤΗΜΙΟ ΘΕΣΣΑΛΟΝΙΚΗΣ

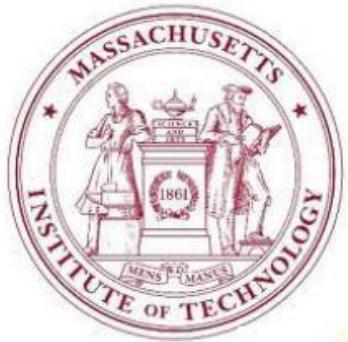
# Introduction

- A full-system computer architecture simulator
  - Open source tool focused on architectural modeling
  - BSD license
- Encompasses
  - system-level architecture, as well as
  - processor micro-architecture
- The gem5 simulation infrastructure is the merger of
  - The best aspects of the M5 and
  - The best aspects of GEMS
- M5
  - Highly configurable simulation framework to support multiple ISAs, and diverse CPU models
  - developed @ The University of Michigan
- GEMS [General Execution-driven Multiprocessor Simulator]
  - detailed and flexible memory system model
  - Includes support for **multiple cache coherence protocols** and interconnect models
  - developed @ The University of Wisconsin Madison



# Where did it come from

---



ARM®

AMD

hp®

i n v e n t

gem5

# Users and contributors

- Widely used in academia and industry
- Contributions from
  - ARM, AMD, Google, ...
  - Wisconsin, Cambridge, Michigan, BSC, ...

## In a Nutshell, gem5...

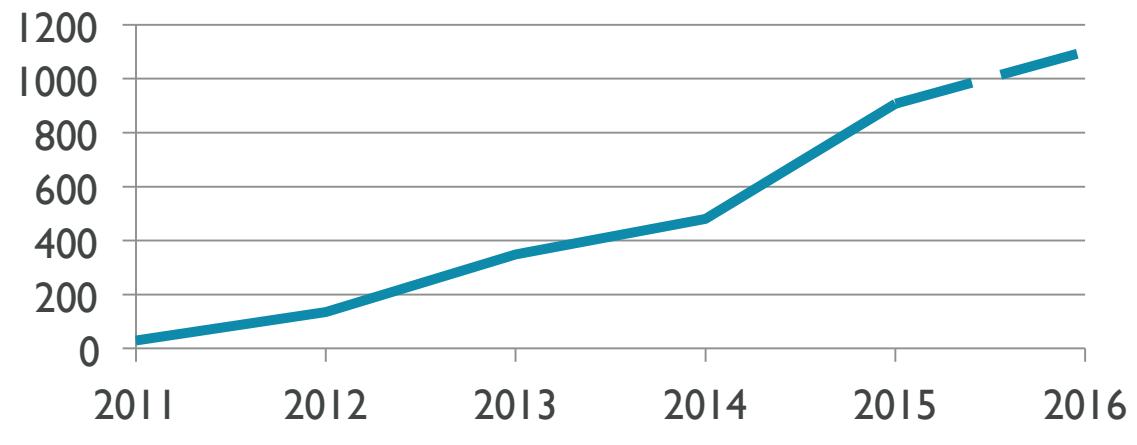
... has had 11,558 commits made by 193 contributors representing 386,321 lines of code

... is mostly written in C++ with a well-commented source code

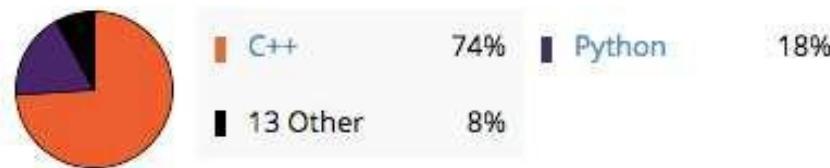
... has a well established, mature codebase maintained by a very large development team with stable Y-O-Y commits

... took an estimated 104 years of effort (COCOMO model) starting with its first commit in October, 2003 ending with its most recent commit 14 days ago

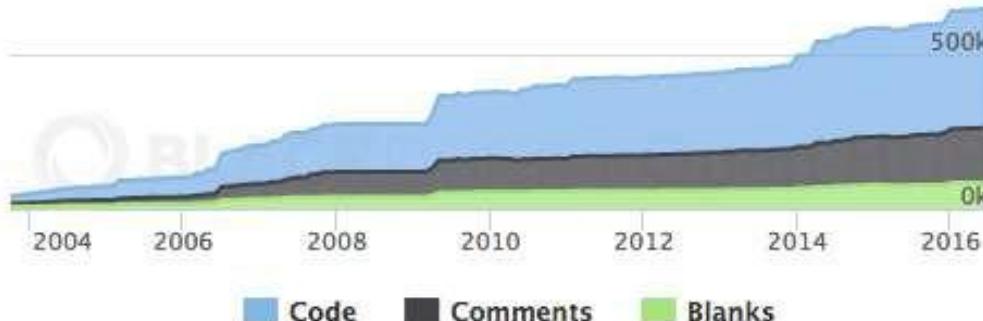
## Publications with gem5



## Languages



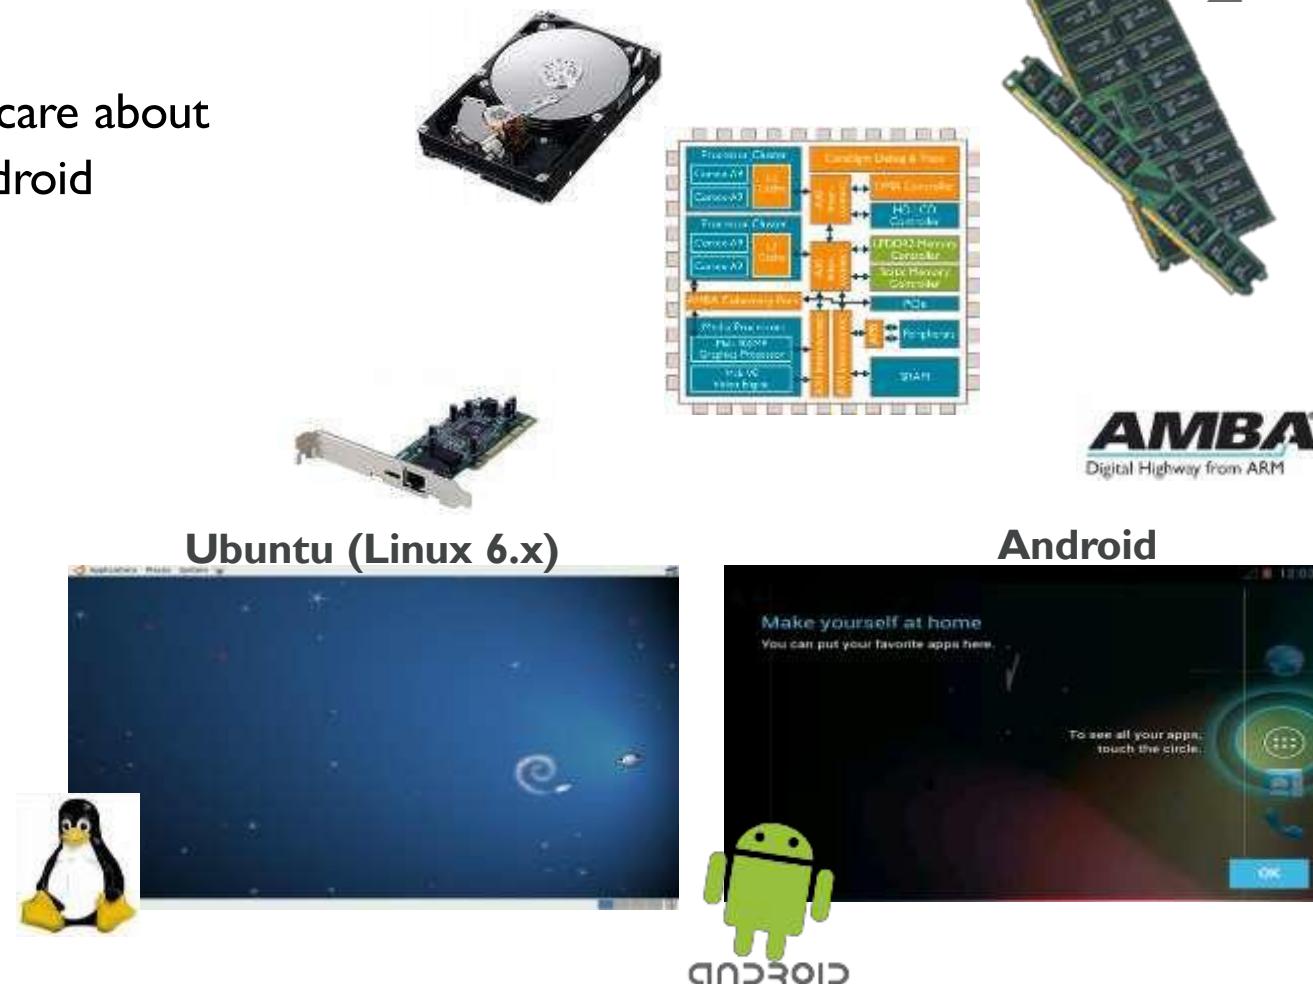
## Lines of Code



# Why gem5?

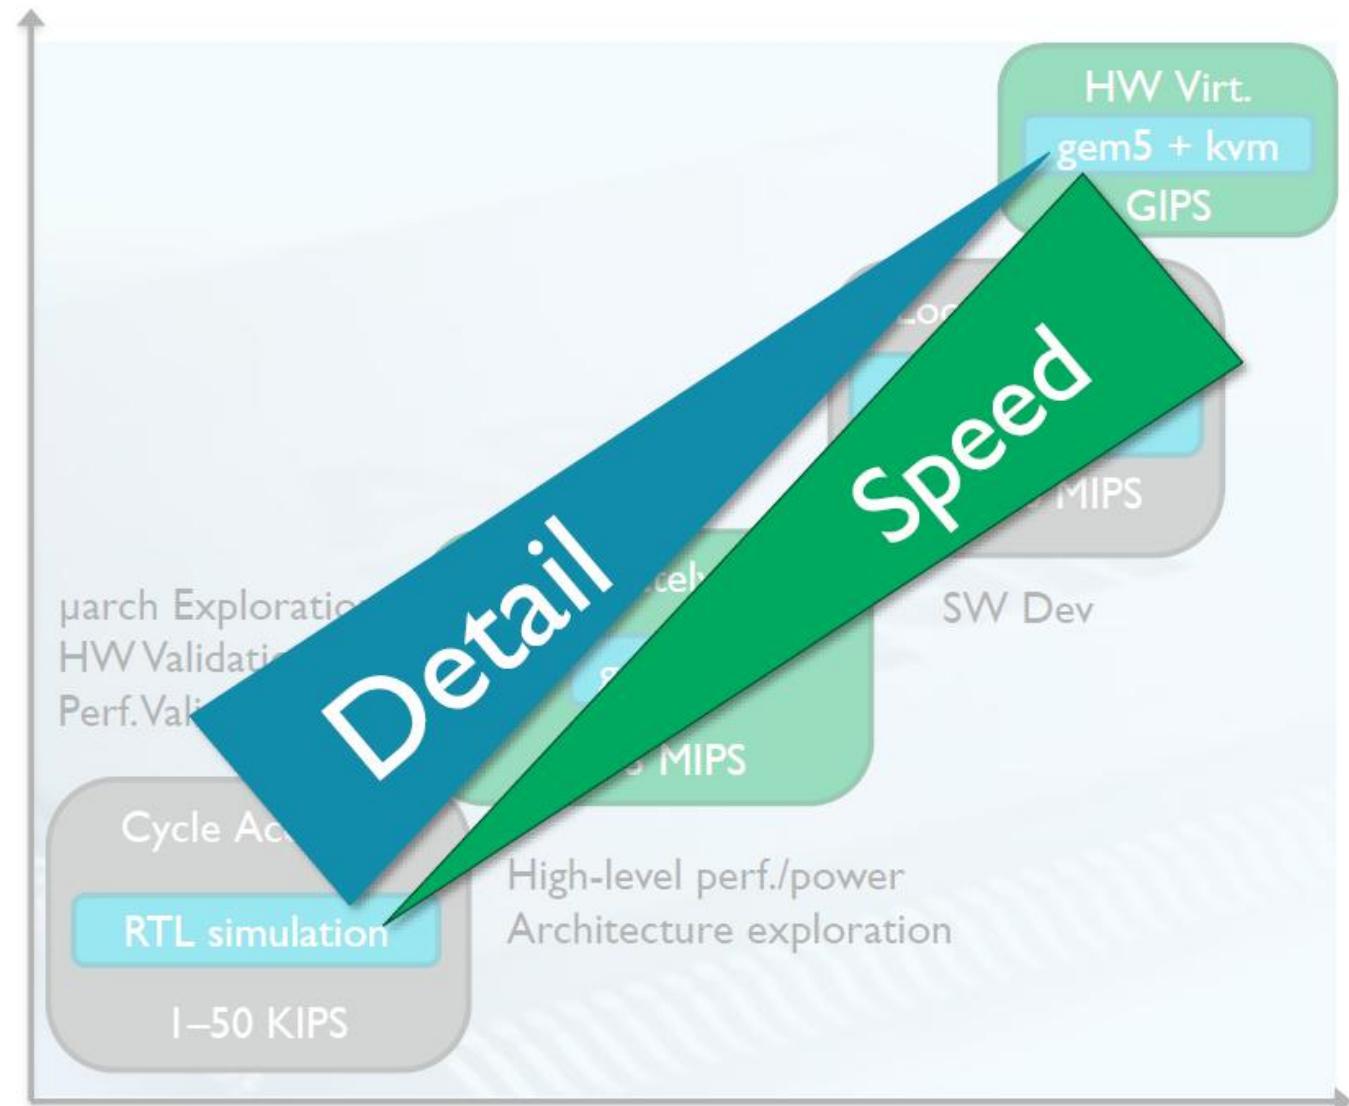


- Runs real workloads
  - Analyze workloads that customers use and care about
  - ... including complex workloads such as Android
- Comprehensive model library
  - Memory and I/O devices
  - Full OS
  - Clients and servers
- Rapid early prototyping
  - New ideas can be tested quickly
  - System-level impact can be quantified
- System-level insights
  - Enables us to study complex memory-system interactions
- Can be wired to custom models
  - Add detail where it matters, when it matters!



# Level of detail

- HW Virtualization
  - Very no/limited timing
  - The same Host/guest ISA
- Functional mode
  - No timing, chain basic blocks of instructions
  - Can add cache models for warming
- Timing mode
  - Single time for execute and memory lookup
  - Advanced on bundle
- Detailed mode
  - Full out-of-order, in-order CPU models
  - Hit-under-miss, reordering,...

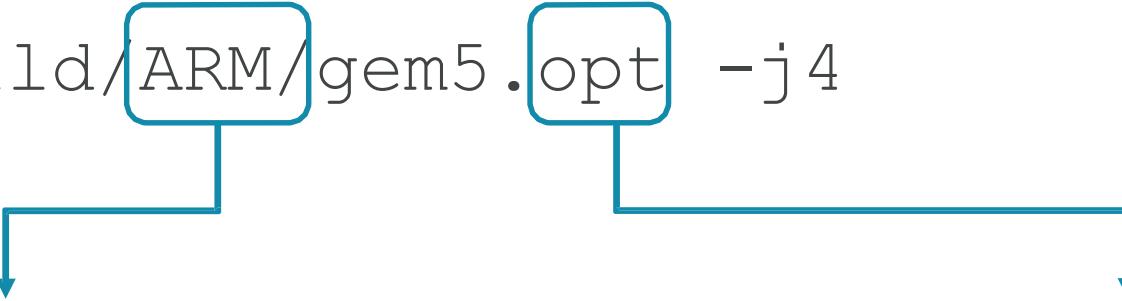


# Prerequisites

- Operating system:
  - OSX, Linux (recommended)
  - Limited support for Windows 10 with a Linux environment
- Software:
  - git
  - Python (dev packages)
  - SCons
  - gcc 4.8 or clang 3.1 (or newer)
  - SWIG 2.0.4 or newer
  - make
- Optional:
  - dtc (to compile device trees)
  - ARMv8 cross compilers (to compile workloads)
  - python-pydot (to generate system diagrams)

# Compiling gem5

```
$ scons build/ARM/gem5.opt -j 4
```



- Guest architecture
- Several architectures in the source tree.
- Most common ones are:
  - **ARM**
  - **RISC-V**
  - **X86**
- Optimization level:
  - **debug**: Debug symbols, no/few optimizations
  - **opt**: Debug symbols + most optimizations
  - **fast**: No symbols + even more optimizations

# Modes

- gem5 has two fundamental modes
- Full system (FS)
  - For **booting operating systems (OSs)**
  - Including devices
  - Interrupts, exceptions, privileged instructions
- Syscall emulation (SE)
  - For running individual applications, or set of applications on MP
  - Models user-visible ISA plus common system calls
  - System calls emulated, typ. by calling host OS
  - Simplified address translation model, no scheduling

# Capabilities

- **Execution modes:** System-call Emulation (SE) & Full-System (FS)
- **ISAs:** Alpha, **ARM**, **RISC-V**, MIPS, Power, SPARC, **x86**
- **CPU models:** AtomicSimple, TimingSimple, InOrder, and O3
- **Cache coherence protocols:** broadcast-based, directories, etc.
- **Interconnection networks:** Simple & Garnet (Princeton, MIT)
- **Devices:** NICs, IDE controller, etc.
- **Multiple systems:** communicate over TCP/IP

# Flexibility

Processor		Memory System		
CPU Model	System Mode	Classic	Ruby	
			Simple	Garnet
Atomic	SE			
	FS	Speed		
Timing	SE			
	FS			
In-Order	SE			
	FS			
O3	SE			
	FS			



The diagram illustrates a spectrum of system configurations based on Speed and Accuracy. The 'Speed' column is represented by the 'Classic' memory system, while the 'Accuracy' column is represented by the 'Ruby' memory system. The arrow indicates that as speed increases, accuracy tends to decrease.

Figure 1: Speed vs. Accuracy Spectrum.

- Classic (from M5): Fast and configurable memory system model.
- Ruby (from GEMS) : framework/infrastructure to model variety of cache-coherent memory system.

Source: [The gem5 Simulator](#), May 2011 issue of ACM SIGARCH Computer Architecture News

# Example disk images

- Example kernels and disk images can be downloaded from gem5.org/Download
  - This includes pre-compiled boot loaders
  - Old but useful to get started
- Download and extract this into a new directory:
  - wget <http://www.gem5.org/dist/current/arm/aarch-system-2014-10.tar.xz>
  - mkdir dist; cd dist
  - tar xvf ../aarch-system-2014-10.tar.xz
- Set the M5\_PATH variable to point to this directory:
  - export M5\_PATH=/path/to/dist
- Most example scripts try to find files using M5\_PATH
  - Kernels/boot loaders/device trees in \${M5\_PATH}/binaries
  - Disk images in \${M5\_PATH}/disks

# Compiling Linux for gem5

```
1. sudo apt install gcc-aarch64-linux-gnu  
2. git clone -b gem5/v4.4 https://github.com/gem5/linux-arm-gem5  
3. cd linux-arm-gem5  
4. make ARCH=arm64 CROSS_COMPILE=aarch64-linux-gnu- gem5_defconfig  
5. make ARCH=arm64 CROSS_COMPILE=aarch64-linux-gnu- -j `nproc`
```

- Builds the default kernel configuration for gem5
  - Has support for most of the devices that gem5 supports

# Running an example script

```
$ build/ARM/gem5.opt configs/example/arm/fs_bigLITTLE.py \
  --kernel path/to/vmlinux \
  --cpu-type atomic \
  --dtb $PWD/system/arm/dt/armv8_gem5_v1_big_little_l_l.dtb \
  --disk your_disk_image.img
```

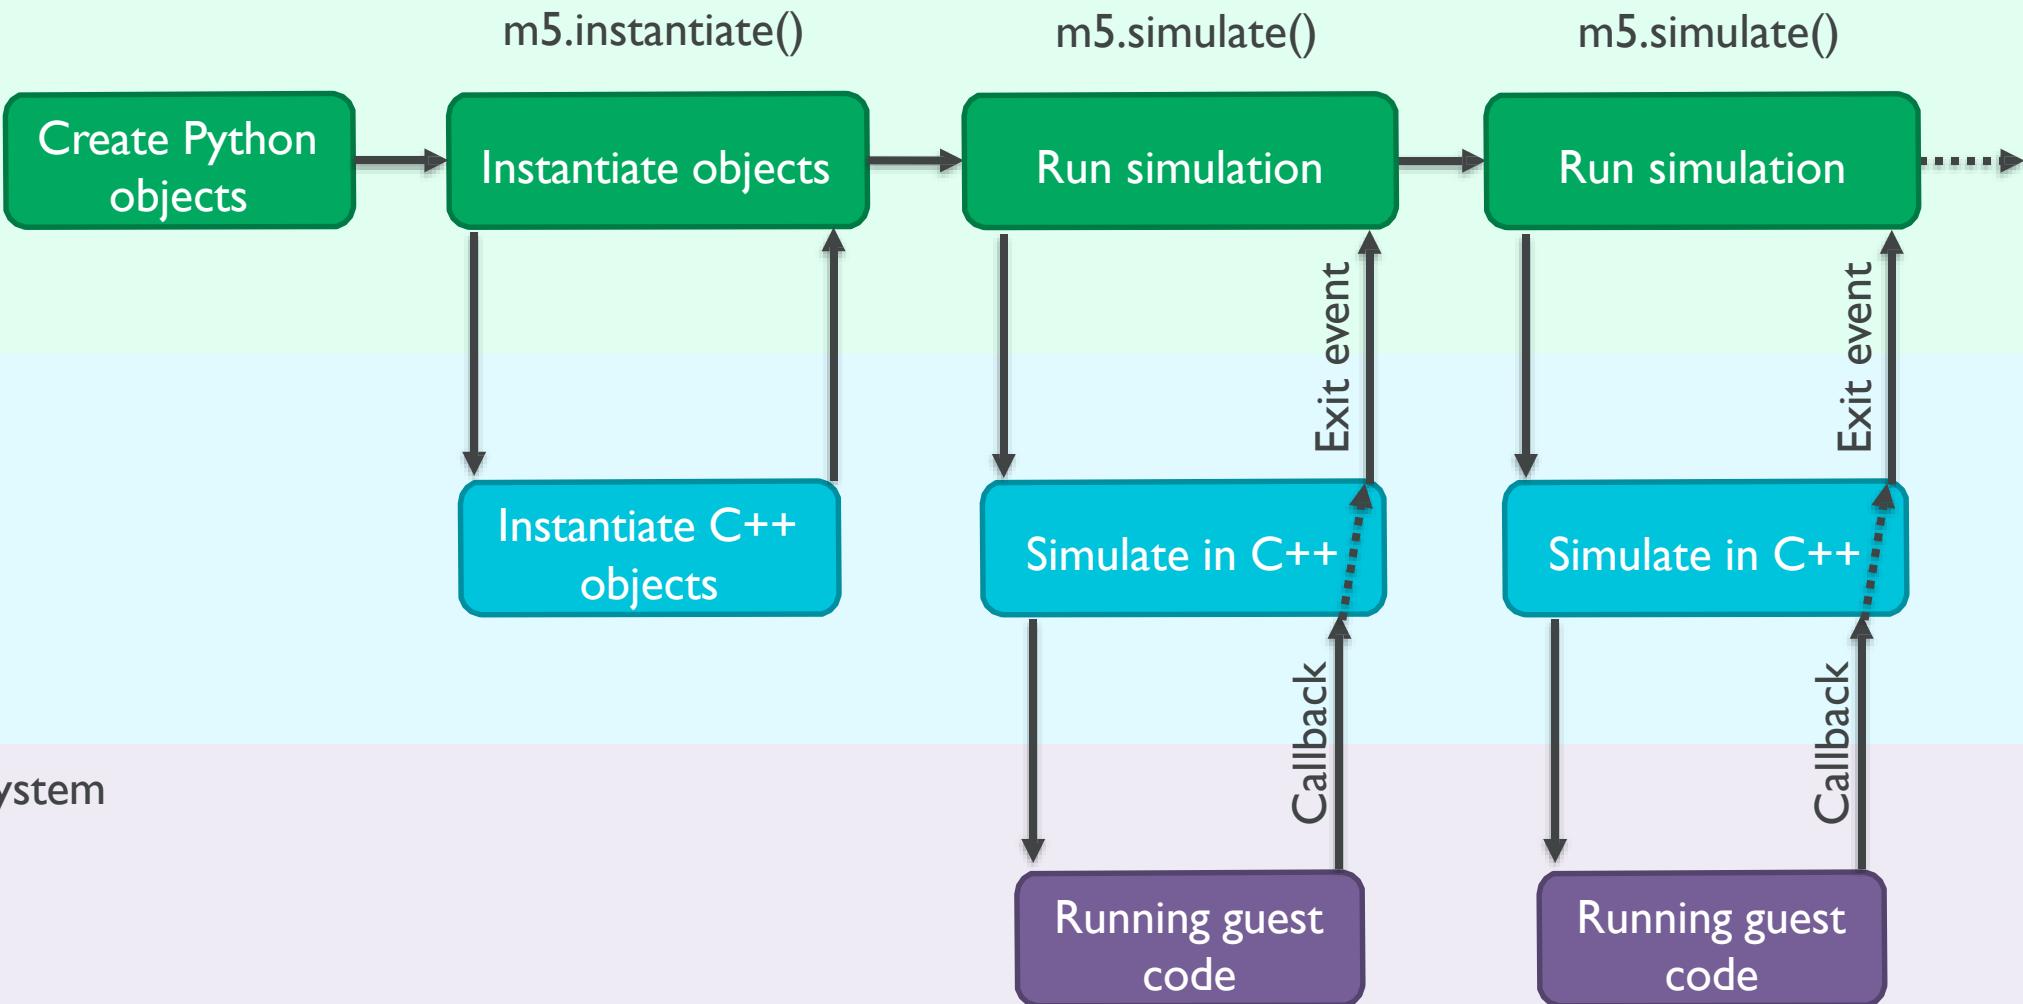
- Simulates a bL system with 1+1 cores
  - Uses a functional ‘atomic’ CPU model
  - Use the ‘timing’ CPU type for an example OoO + InO configuration

# Design philosophy

- gem5 is conceptually a Python library implemented in C++
  - **Configured** by instantiating Python classes with matching C++ classes
  - **Model parameters** exposed as attributes in Python
  - **Running** is controlled from Python, but implemented in C++
- Configuration and running are two distinct steps
  - **Configuration phase** ends with a call to instantiate the C++ world
  - **Parameters cannot be changed** after the C++ world has been created

# Control flow

Python



# Overriding model parameters

```
import m5
```

```
class L1DCache(m5.objects.Cache):  
    assoc = 2  
    size = '16kB'
```

- Use gem5's base Cache
- Override associativity
- Override size

```
class L1ICache(L1DCache):  
    assoc = 16
```

- Use defaults from L1DCache
- Override associativity again

```
l1i = L1ICache(assoc=8,  
               repl=m5.objects.RandomRepl())
```

- Override parameters at instantiation time
- We'll cover memory ports later

# Running

```
m5.instantiate()
```

- Instantiate the C++ world

```
event = m5.simulate()
```

- Start the simulation

```
print 'Exiting @ tick %i: %s' \
      % ( m5.curTick(),
          event.getCause() )
```

- Print why the simulator exited
- Sometimes desirable to call  
m5.simulate() again.

```
m5.simulate(m5.tick.fromSeconds(0.1))
```

- Run for a fixed number of  
simulated seconds

# Creating Checkpoints

```
m5.checkpoint('name.cpt')
```

- Checkpoints can be used to **store the simulator's state**
  - Can be used to implement SimPoints or similar methodologies
- Checkpoint limitations:
  - The act of taking a checkpoint affects system state!
  - **Checkpoints don't store cache state**
  - **Checkpoints don't store pipeline state**

# Restoring Checkpoints

```
m5.instantiate('name.cpt')
```

```
event = m5.simulate()
```

- Instantiate system and load state from checkpoint
- Run in the same way as before

# Dumping statistics

- Can be requested from Python:
  - `m5.stats.dump()` : Dump statistics
  - `m5.stats.reset()`: Reset stat counters
- Guest command line:
  - `m5 dumpstats [ [delay] [period] ]`
  - `m5 dumpresetstats [ [delay] [period] ]`
- Guest code using libm5.a:
  - `m5_dump_stats(delay, periodicity)` : Dump statistics
  - `m5_dumpreset_stats(delay, periodicity)` : Dump & reset statistics

# Understanding gem5 output

```
> ls m5out
```

[config.ini](#)

[config.json](#)

[stats.txt](#)

**config.ini:** Dumps all of the parameters of all SimObjects (modules). This shows exactly what you simulated.

**config.json:** Same as config.ini, but in json format.

**stats.txt:** Detailed statistic output. Each SimObject defines and updates statistics. They are printed here at the end of simulation.

# stats.txt

```
----- Begin Simulation Statistics -----  
sim_seconds      0.000346      # Number of seconds simulated  
sim_ticks        345518000     # Number of ticks simulated  
final_tick        345518000     # Number of ticks from start  
sim_freq         1000000000000 # Frequency of simulation  
...  
sim_insts        5712          # Number of instructions simulated  
sim_ops           10314         # Number of ops (including micro)  
...  
system.mem_ctrl.bytes_read::cpu.inst 58264 # Number of instruction reads  
system.mem_ctrl.bytes_read::cpu.data 7167 # Number of data reads  
...  
system.cpu.committedOps            10314 # Number of committed ops  
system.cpu.num_int_alu_accesses 10205 # Number of integer ALU accesses
```

**sim\_seconds:** name of stat. This shows *simulated guest time*

Every SimObject can have its own stats. Names are what you used in the Python config file

# Debugging

# Debugging Facilities

- Tracing
  - Instruction tracing
  - Diffing traces
- Using gdb to debug gem5
  - Debugging C++ and gdb-callable functions
  - Remote debugging
- Pipeline viewer

# Tracing/Debugging

- `printf()` is a nice debugging tool
  - Keep good print statements in code and selectively enable them
  - Lots of debug output can be a very good thing when a problem arises
  - Use `DPRINTF`s in code
  - `DPRINTF(TLB, "Inserting entry into TLB with pfn:%#x...")`
- Example flags:
  - Fetch, Decode, Ethernet, Exec, TLB, DMA, Bus, Cache, O3CPUAll
  - Print out all flags with `./build/ARM/gem5.opt -- debug-help`
- Enabled on the command line
  - `--debug-flags=Exec`
  - `--debug-start=30000`
  - `--debug-file=my_trace.out`
  - Enable the flag Exec; Start at tick 30000; Write to `my_trace.out`

# Sample Run with Debugging

Command Line:

```
22:44:28 [/work/gem5] ./build/ARM/gem5.opt --debug-flags=Decode --  
debug-start=50000-- debug-file=my_trace.out configs/example/se.py -c  
tests/test-progs/hello/bin/arm/linux/hello  
...  
***** REAL SIMULATION *****  
info: Entering event queue @ 0. Starting simulation...  
Hello world!  
Exiting @ tick 3107500 because target called exit()
```

**my\_trace.out:**

```
2:44:47 [ /work/gem5] head m5out/my_trace.out 0xe353001e  
50000: system.cpu: Decode: Decoded cmps instruction:  
50500: system.cpu: Decode: Decoded ldr instruction: 0x979ff103  
51000: system.cpu: Decode: Decoded ldr instruction: 0xe5107004  
51500: system.cpu: Decode: Decoded ldr instruction: 0xe4903008  
52000: system.cpu: Decode: Decoded addi_uop instruction: 0xe4903008  
52500: system.cpu: Decode: Decoded cmps instruction: 0xe3530000  
53000: system.cpu: Decode: Decoded b instruction: 0x1affff84  
53500: system.cpu: Decode: Decoded sub instruction: 0xe2433003  
54000: system.cpu: Decode: Decoded cmps instruction: 0xe353001e  
54500: system.cpu: Decode: Decoded ldr instruction: 0x979ff103
```

# Adding Your Own Flag

- Print statements put in source code
  - Encourage you to add ones to your models or contribute ones you find particularly useful
- Macros remove them from the `gem5.fast` binary
  - There is no performance penalty for adding them
  - To enable them you need to run `gem5.opt` or `gem5.debug`
- Adding one with an existing flag
  - `DPRINTF(<flag>, "normal printf %s\n", "arguments");`
- To add a new flag add the following in a Sconscript
  - `DebugFlag('MyNewFlag')`
  - Include corresponding header, e.g. `#include "debug/MyNewFlag.hh"`

Ευχαριστώ για την προσοχή σας!  
απορίες?



ΤΜΗΜΑ ΗΛΕΚΤΡΟΛΟΓΩΝ ΜΗΧΑΝΙΚΩΝ  
& ΜΗΧΑΝΙΚΩΝ ΥΠΟΛΟΓΙΣΤΩΝ

ΑΡΙΣΤΟΤΕΛΕΙΟ ΠΑΝΕΠΙΣΤΗΜΙΟ ΘΕΣΣΑΛΟΝΙΚΗΣ